
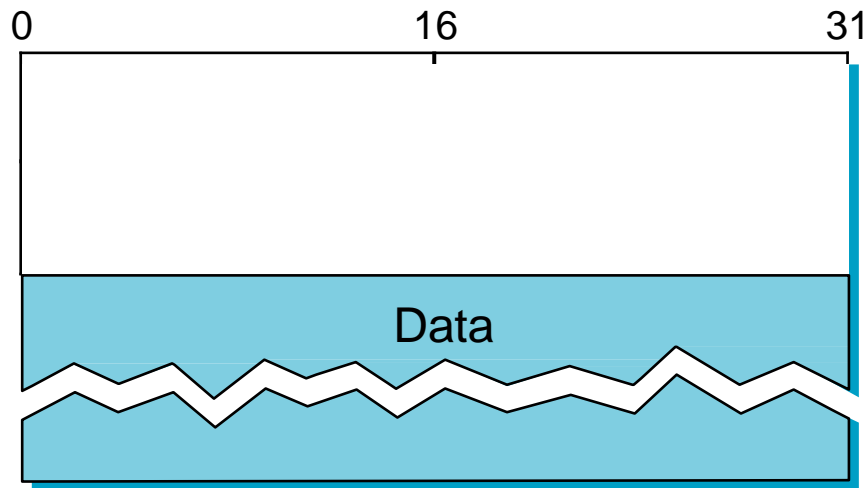


Socket Programming

Demultiplexing

- Convert host-to-host packet delivery service into a process-to-process communication channel



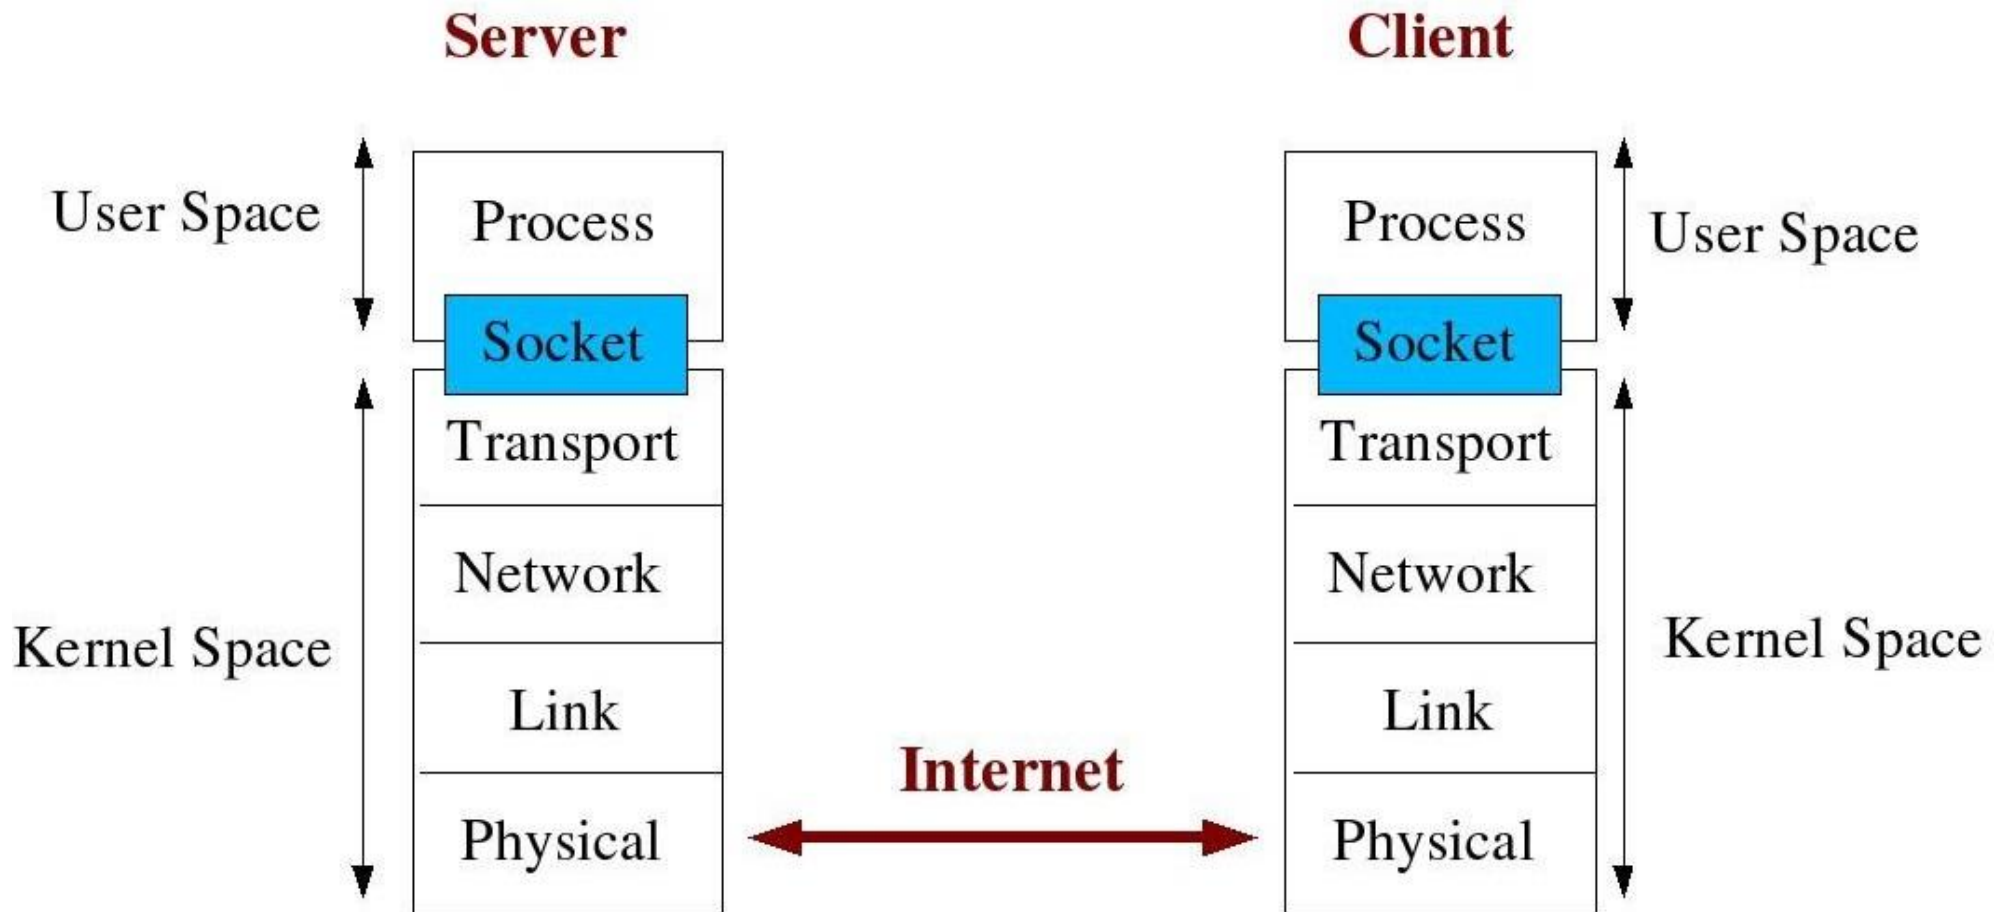
Byte Ordering

- Two types of “Byte ordering”
 - Network Byte Order: High-order byte of the number is stored in memory at the lowest address
 - Host Byte Order: Low-order byte of the number is stored in memory at the lowest address
 - Network stack (TCP/IP) expects Network Byte Order
- Conversions:
 - htons() - Host to Network Short
 - htonl() - Host to Network Long
 - ntohs() - Network to Host Short
 - ntohl() - Network to Host Long

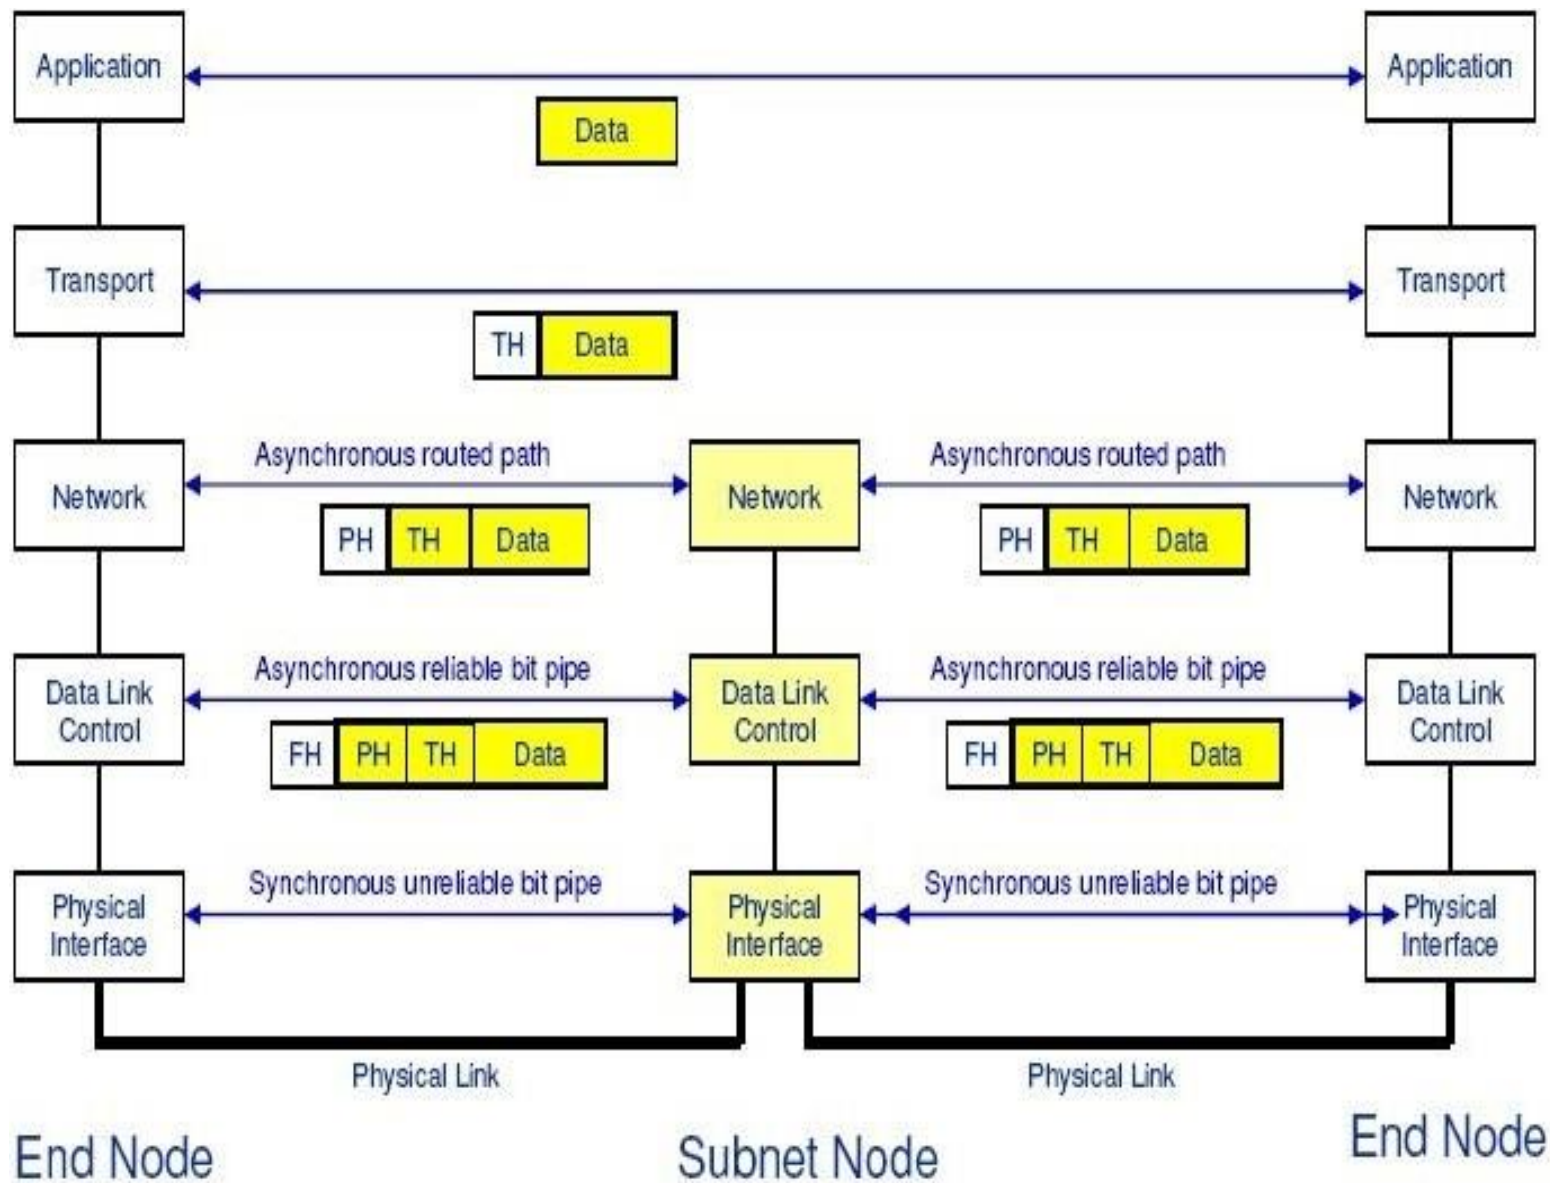
What is a socket?

- Socket: An interface between an application process and transport layer
 - The application process can send/receive messages to/from another application process (local or remote) via a socket
- In Unix jargon, a socket is a file descriptor – an integer associated with an open file
- Types of Sockets: **Internet Sockets**, unix sockets, X.25 sockets etc
 - Internet sockets characterized by IP Address (4 bytes), port number (2 bytes)

Socket Description



Encapsulation



Each layer just looks at its own header

Types of Internet Sockets

- Stream Sockets (SOCK_STREAM)
 - Connection oriented
 - Rely on TCP to provide reliable two-way connected communication
- Datagram Sockets (SOCK_DGRAM)
 - Rely on UDP
 - Connection is unreliable

socket() -- Get the file descriptor

- `int socket(int domain, int type, int protocol);`
 - domain should be set to `PF_INET`
 - type can be `SOCK_STREAM` or `SOCK_DGRAM`
 - set protocol to 0 to have socket choose the correct protocol based on type
 - `socket()` returns a socket descriptor for use in later system calls or -1 on error

```
int sockfd;
```

```
sockfd = socket (PF_INET, SOCK_STREAM, 0);
```


Socket Structures

- struct sockaddr: Holds socket address information for many types of sockets

```
struct sockaddr {  
    unsigned short  sa_family;    //address family AF_XXX  
    unsigned short  sa_data[14]; //14 bytes of protocol addr  
}
```

- struct sockaddr_in: A parallel structure that makes it easy to reference elements of the socket address

```
struct sockaddr_in {  
    short int          sin_family;    // set to AF_INET  
    unsigned short int sin_port;      // Port number  
    struct in_addr     sin_addr;      // Internet address  
    unsigned char       sin_zero[8];  //set to all zeros  
}
```

- sin_port and sin_addr must be in **Network Byte Order**

Dealing with IP Addresses

- ```
struct in_addr {
 unsigned long s_addr; // that's a 32bit long, or 4 bytes
};
```
- ```
int inet_aton(const char *cp, struct in_addr *inp);  
  
struct sockaddr_in  my_addr;  
my_addr.sin_family = AF_INET;  
my_addr.sin_port = htons(MYPORT);  
inet_aton("10.0.0.5",&(my_addr.sin_addr));  
memset(&(my_addr.sin_zero),'\0',8);  
  
- inet_aton() gives non-zero on success; zero on failure
```
- To convert binary IP to string: `inet_ntoa()`

```
printf("%s",inet_ntoa(my_addr.sin_addr));
```

bind() - what port am I on?

- Used to associate a socket with a port on the local machine
 - The port number is used by the kernel to match an incoming packet to a process
- `int bind(int sockfd, struct sockaddr *my_addr, int addrlen)`
 - `sockfd` is the socket descriptor returned by `socket()`
 - `my_addr` is pointer to `struct sockaddr` that contains information about your IP address and port
 - `addrlen` is set to `sizeof(struct sockaddr)`
 - returns -1 on error
 - `my_addr.sin_port = 0;` //choose an unused port at random
 - `my_addr.sin_addr.s_addr = INADDR_ANY;` //use my IP adr

Example

```
int sockfd;  
  
struct sockaddr_in my_addr;  
  
sockfd = socket(PF_INET, SOCK_STREAM, 0);  
  
my_addr.sin_family = AF_INET;           // host byte order  
my_addr.sin_port = htons(MYPORT);      // short, network byte  
                                         order  
  
my_addr.sin_addr.s_addr = inet_addr("172.28.44.57");  
  
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct  
  
bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct  
sockaddr));  
  
/***** Code needs error checking. Don't forget to do that *****/
```

connect() - Hello!

- Connects to a remote host
- `int connect(int sockfd, struct sockaddr *serv_addr, int addrlen)`
 - `sockfd` is the socket descriptor returned by `socket()`
 - `serv_addr` is pointer to `struct sockaddr` that contains information on destination IP address and port
 - `addrlen` is set to `sizeof(struct sockaddr)`
 - returns -1 on error
- No need to `bind()`, kernel will choose a port

Example

```
#define DEST_IP  "172.28.44.57"
#define DEST_PORT 5000
main(){
    int sockfd;
    struct sockaddr_in dest_addr;  // will hold the destination addr
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
    dest_addr.sin_family = AF_INET;        // host byte order
    dest_addr.sin_port = htons(DEST_PORT); // network byte
    order
    dest_addr.sin_addr.s_addr = inet_addr(DEST_IP);
    memset(&(dest_addr.sin_zero), '\0', 8); // zero the rest of the
    struct    connect(sockfd, (struct sockaddr *)&dest_addr,
    sizeof(struct sockaddr));
    /***** Don't forget error checking *****/
}
```

listen() - Call me please!

- Waits for incoming connections
- `int listen(int sockfd, int backlog);`
 - `sockfd` is the socket file descriptor returned by `socket()`
 - `backlog` is the number of connections allowed on the incoming queue
 - `listen()` returns -1 on error
 - Need to call `bind()` before you can `listen()`
 - `socket()`
 - `bind()`
 - `listen()`
 - `accept()`

accept() - Thank you for calling !

- accept() gets the pending connection on the port you are listen()ing on
- `int accept(int sockfd, void *addr, int *addrlen);`
 - sockfd is the listening socket descriptor
 - information about incoming connection is stored in addr which is a pointer to a local struct `sockaddr_in`
 - addrlen is set to `sizeof(struct sockaddr_in)`
 - accept returns *a new socket file descriptor* to use for this accepted connection and -1 on error

Example

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#define MYPORT 3490    // the port users will be connecting to
#define BACKLOG 10    // pending connections queue will hold
main(){
    int sockfd, new_fd; // listen on sock_fd, new connection on
    new_fd
    struct sockaddr_in my_addr;    // my address information
    struct sockaddr_in their_addr; // connector's address information
    int sin_size;
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
```

Cont...

```
my_addr.sin_family = AF_INET;          // host byte order
my_addr.sin_port = htons(MYPORT);      // short, network byte
order
my_addr.sin_addr.s_addr = INADDR_ANY; // auto-fill with my IP
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct

// don't forget your error checking for these calls:

bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct
sockaddr));

listen(sockfd, BACKLOG);

sin_size = sizeof(struct sockaddr_in);

new_fd = accept(sockfd, (struct sockaddr *)&their_addr,
&sin_size);
```

send() and recv() - Let's talk!

- The two functions are for communicating over stream sockets or connected datagram sockets.
- `int send(int sockfd, const void *msg, int len, int flags);`
 - `sockfd` is the socket descriptor you want to send data to (returned by `socket()` or got from `accept()`)
 - `msg` is a pointer to the data you want to send
 - `len` is the length of that data in bytes
 - set `flags` to 0 for now
 - `send()` returns the number of bytes actually sent (may be less than the number you told it to send) or -1 on error

send() and recv() - Let's talk!

- `int recv(int sockfd, void *buf, int len, int flags);`
 - `sockfd` is the socket descriptor to read from
 - `buf` is the buffer to read the information into
 - `len` is the maximum length of the buffer
 - set `flags` to 0 for now
 - `recv()` returns the number of bytes actually read into the buffer or -1 on error
 - If `recv()` returns 0, the remote side has closed connection on you

sendto() and recvfrom() - DGRAM style

- `int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);`
 - *to* is a pointer to a struct `sockaddr` which contains the destination IP and port
 - *tolen* is `sizeof(struct sockaddr)`
- `int recvfrom(int sockfd, void *buf, int len, int flags, struct sockaddr *from, int *fromlen);`
 - *from* is a pointer to a local struct `sockaddr` that will be filled with IP address and port of the originating machine
 - *fromlen* will contain length of address stored in *from*

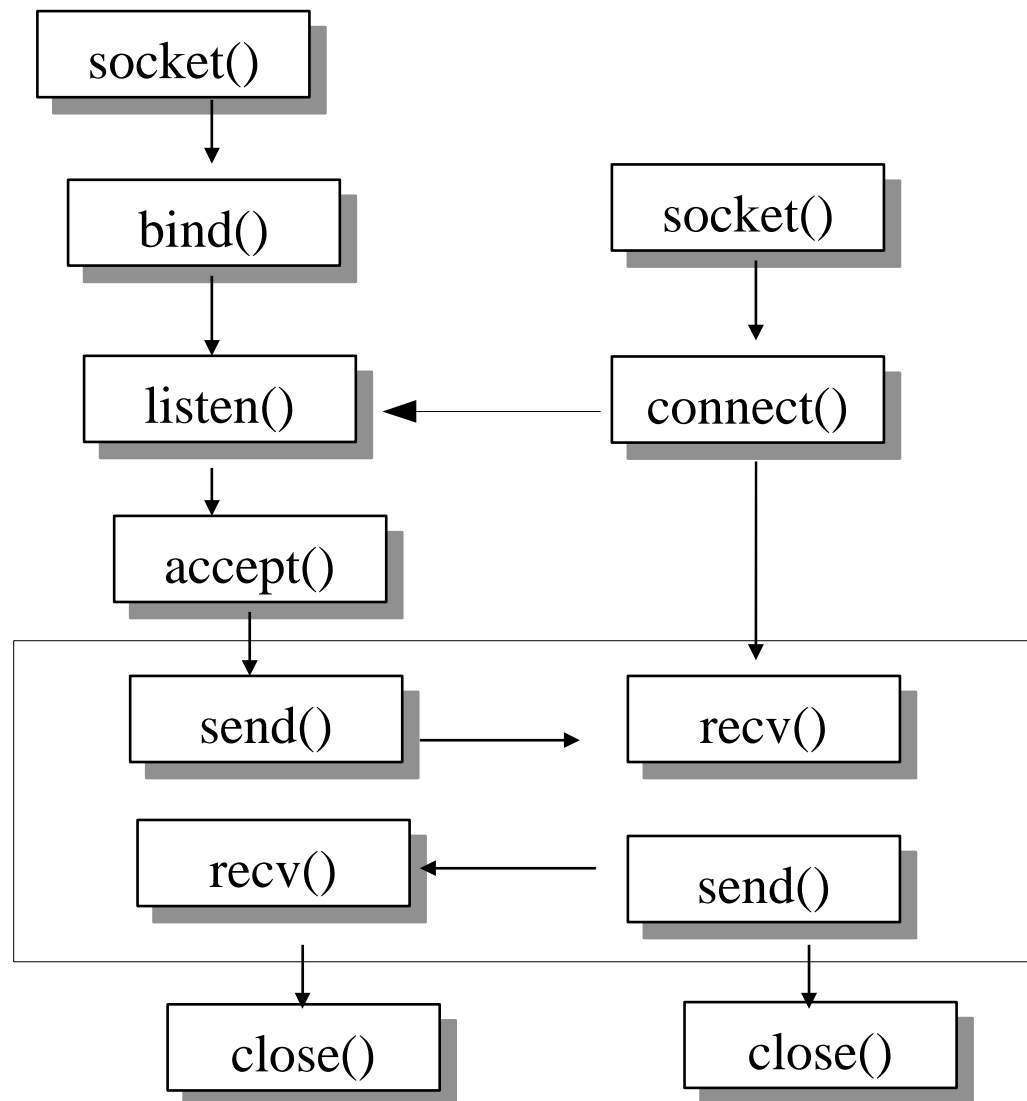
close() - Bye Bye!

- `int close(int sockfd);`
 - Closes connection corresponding to the socket descriptor and frees the socket descriptor
 - Will prevent any more sends and recvs

Connection Oriented Protocol

Server

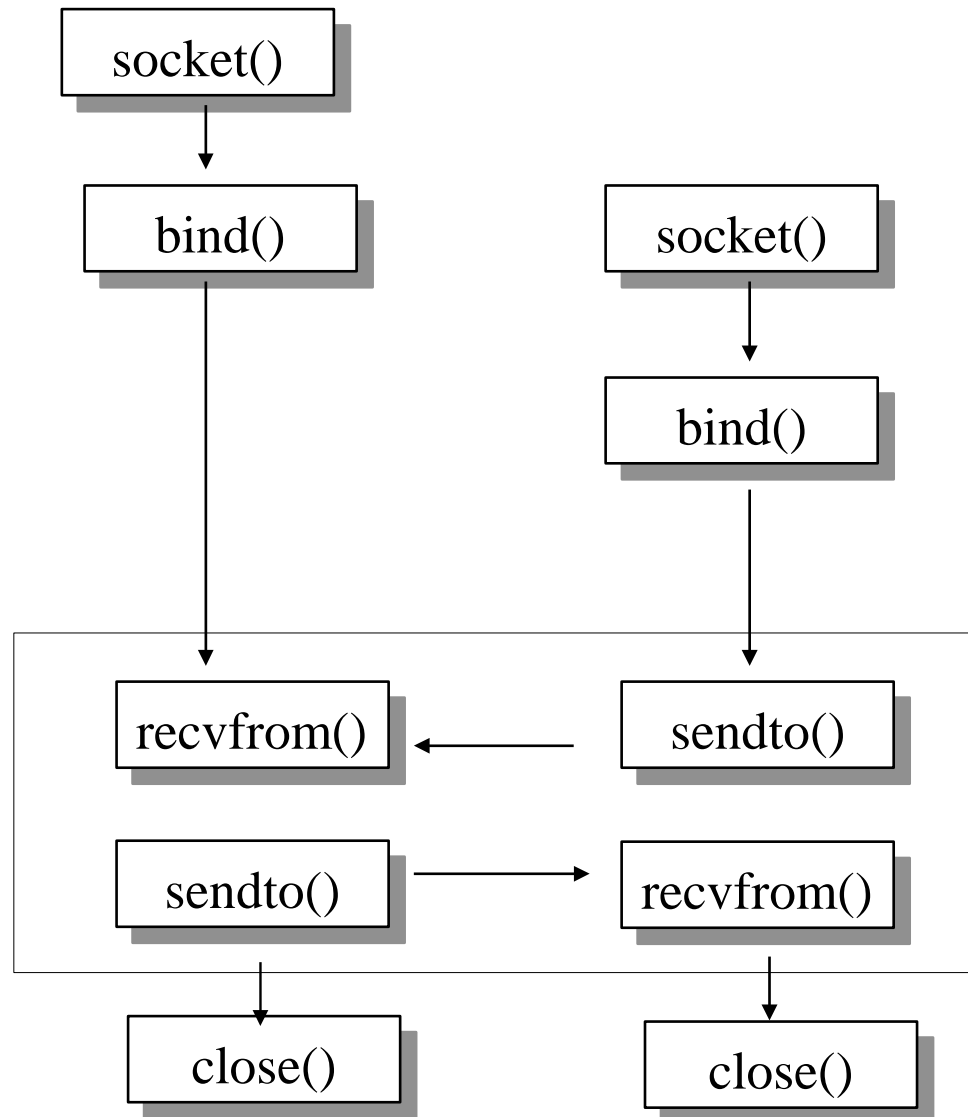
Client



Connectionless Protocol

Server

Client



Miscellaneous Routines

- `int getpeername(int sockfd, struct sockaddr *addr, int *addrlen);`
 - Will tell who is at the other end of a connected stream socket and store that info in *addr*
- `int gethostname(char *hostname, size_t size);`
 - Will get the name of the computer your program is running on and store that info in *hostname*

Miscellaneous Routines

- `struct hostent *gethostbyname(const char *name);`
 `struct hostent {`
 `char *h_name; //official name of host`
 `char **h_aliases; //alternate names for the host`
 `int h_addrtype; //usually AF_NET`
 `int h_length; //length of the address in bytes`
 `char **h_addr_list; //array of network addresses for the host`
 `}`
 `#define h_addr h_addr_list[0]`

- **Example Usage:**

```
struct hostent *h;  
h = gethostbyname("www.iitk.ac.in");  
printf("Host name : %s \n", h->h_name);  
printf("IP Address: %s\n",inet_ntoa(*((struct in_addr *)h->h_addr)));
```

Advanced Topics

- Blocking
- Select
- Handling partial sends
- Signal handlers
- Threading

Summary

- Sockets help application process to communicate with each other using standard Unix file descriptors
- Two types of Internet sockets: `SOCK_STREAM` and `SOCK_DGRAM`
- Many routines exist to help ease the process of communication

References

- Books:
 - Unix Network Programming, volumes 1-2 by W. Richard Stevens.
 - TCP/IP Illustrated, volumes 1-3 by W. Richard Stevens and Gary R. Wright
- Web Resources:
 - Beej's Guide to Network Programming
 - www.ecst.csuchico.edu/~beej/guide/net/